

IE417 : Embedded AI

Lab 1 - Voice Controlling LEDs with Edge Impulse

Submitted by

Avinash Baraiya- 202201211

Malhar Vaghasiya - 202201183

Smeet Agrawal - 202101237

Priyesh Tandel - 202101222

Samarth Panchal- 202101456

Submitted to

Professor Tapas Kumar Maiti



**Dhirubhai Ambani Institute of Information and
Communication Technology**
Gandhinagar , Gujarat

Goal

In this lab, you are going to implement the following recipes:

1. Acquiring audio data with a laptop or smartphone
2. Extracting MFCC features from audio samples
3. Designing and training a neural network (NN) model
4. Tuning model performance with EON Tuner
5. Live classifications with a smartphone
6. Live classifications with the Arduino Nano
7. Continuous inferencing on the Arduino Nano

Video Link

Click below for YouTube Demo video



or Paste link in browser - <https://youtu.be/Hpj7VdscqR0>

Screenshots

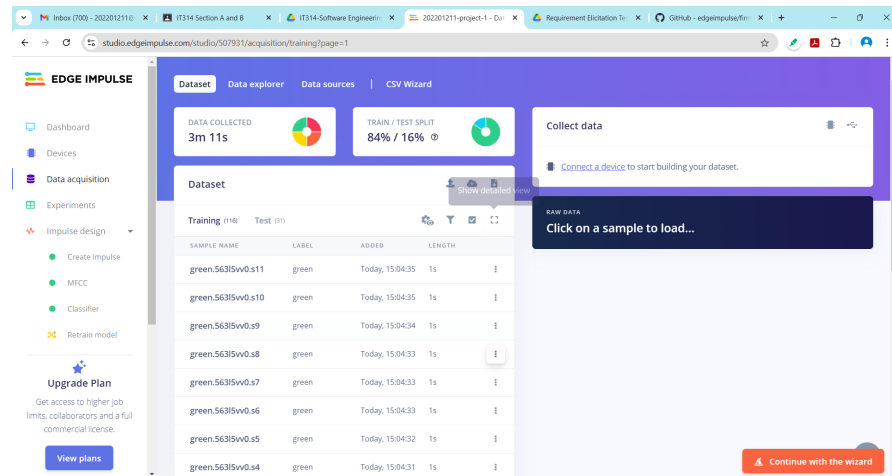


FIGURE 1: Training

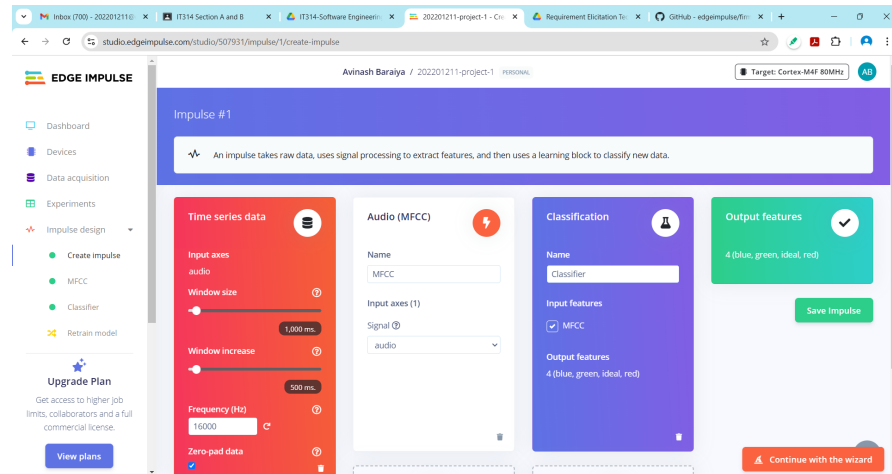


FIGURE 2: Create Impulse

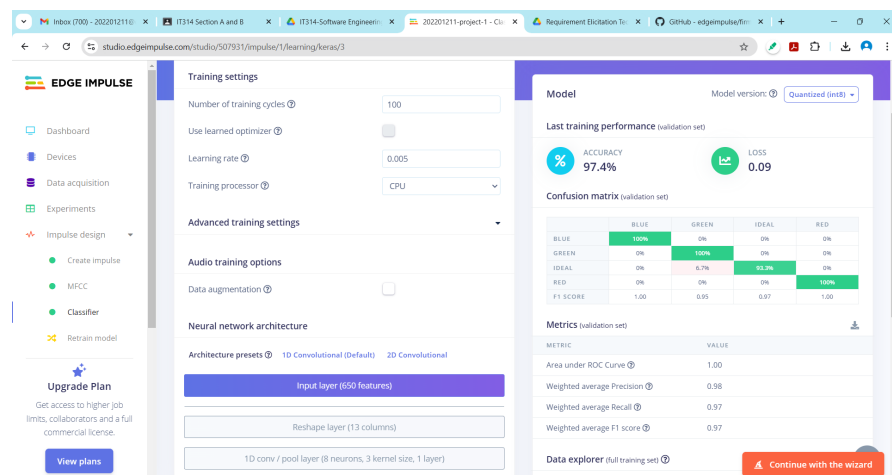


FIGURE 3: Classifier

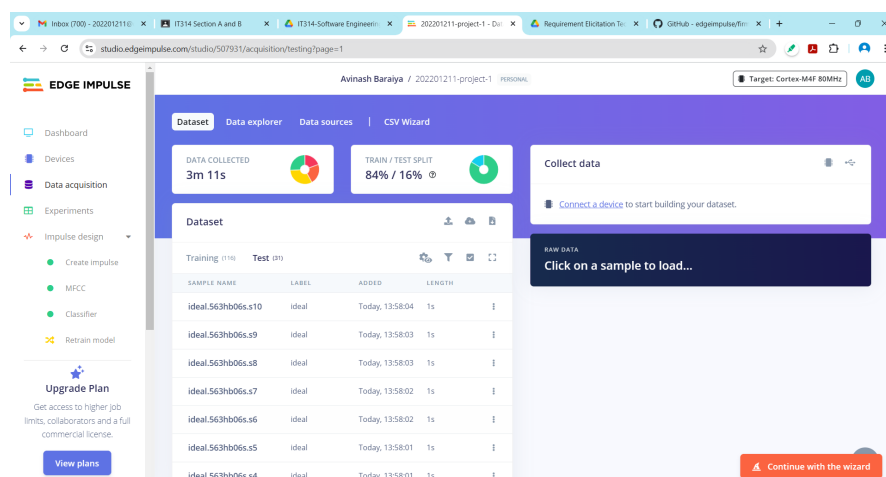


FIGURE 4: Data Acquisition

```

1 #define EIDSP_QUANTIZE_FILTERBANK 0
2
3 /**
4  * Define the number of slices per model window. E.g. a model window of 1000 ms
5  * with slices per model window set to 4, results in a slice size of 250 ms.
6  * For more info: https://docs.edgeimpulse.com/docs/continuous-audio-sampling
7  */
8 #define EI_CLASSIFIER_SLICES_PER_MODEL_WINDOW 3
9
10 /* Include
11 #include <Arduino.h>
12 #include <a202201211-project-1_inferencing.h>
13
14 #define RED 22
15 #define BLUE 24
16 #define GREEN 26
17 #define LED_PWM 25
18
19 /** Audio buffers, pointers and selectors */
20 typedef struct {
21     signed short *buffers[3];
22     unsigned char buf_select;
23     unsigned char buf_ready;
24     unsigned int buf_count;
25     unsigned int n_samples;
26 } inference_t;
27
28 static inference_t inference;
29 static bool record_ready = false;
30 static signed short *sampleBuffer;
31 static bool debug_on = false; // Set this to true to see e.g. features generated from the raw signal
32 static int print_results = (EI_CLASSIFIER_SLICES_PER_MODEL_WINDOW);
33
34 /**
35  * @brief Arduino setup function
36  */
37 void setup()
38 {
39     pinMode(RED, OUTPUT);
40     pinMode(BLUE, OUTPUT);
41     pinMode(GREEN, OUTPUT);
42     // pinMode(LED_PWM, OUTPUT);
43
44     // put your setup code here, to run once:
45     Serial.begin(115200);
46
47     Serial.println("Edge Impulse Inferencing Demo");
48
49     // summary of inferencing settings (from model_metadata.h)
50     ei_printf("Inferencing settings:\n");
51     ei_printf("Interval: %.2f ms\n", (float)EI_CLASSIFIER_INTERVAL_MS);
52     ei_printf("Stream size: %d\n", EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE);
53     ei_printf("Sample length: %d ms\n", EI_CLASSIFIER_RAW_SAMPLE_COUNT / 16);
54     ei_printf("No. of classes: %d\n", sizeof(ei_classifier_inferencing_categories) /
55             sizeof(ei_classifier_inferencing_categories[0]));
56
57     run_classifier_init();
58     if (microphone_inference_start(EI_CLASSIFIER_SLICE_SIZE) == false) {
59         ei_printf("ERR: Failed to setup audio sampling\n");
60         return;
61     }
62 }
63
64 /**
65  * @brief Arduino main function. Runs the inferencing loop.
66  */
67 void loop()
68 {
69     bool r = microphone_inference_record();
70     if (!r) {
71         ei_printf("ERR: Failed to record audio ... \n");
72         return;
73     }
74
75     signal_t signal;
76     signal.total_length = EI_CLASSIFIER_SLICE_SIZE;
77     signal.get_data = microphone_audio_signal_get_data;
78     ei_impulse_result_t result = {0};
79
80     EI_IMPULSE_ERROR r = run_classifier_continuous(&signal, &result, debug_on);
81     if (r == EI_IMPULSE_OK) {
82         ei_printf("ERR: Failed to run classifier (%d)\n", r);
83         return;
84     }
85
86     if (result.classification[0].value >= 0.7) {
87         digitalWrite(BLUE, LOW);
88         digitalWrite(GREEN, HIGH);
89         digitalWrite(RED, HIGH);
90     }
91     else if (result.classification[1].value >= 0.7) {
92         digitalWrite(BLUE, HIGH);
93         digitalWrite(GREEN, LOW);
94         digitalWrite(RED, HIGH);
95     }
96     else if (result.classification[2].value >= 0.7) {
97         digitalWrite(BLUE, HIGH);
98         digitalWrite(GREEN, HIGH);
99         digitalWrite(RED, HIGH);
100     }
101     else if (result.classification[3].value >= 0.7) {
102         digitalWrite(BLUE, HIGH);
103         digitalWrite(GREEN, HIGH);
104         digitalWrite(RED, LOW);
105     }
106
107     if (print_results >= (EI_CLASSIFIER_SLICES_PER_MODEL_WINDOW)) {
108         // print the predictions
109         ei_printf("Predictions: ");
110         ei_printf("DSP: %d ms, Classification: %d ms, Anomaly: %d ms\n",
111             result.timing.dsp, result.timing.classification, result.timing.anomaly);
112         ei_printf("\n");
113         for (size_t ix = 0; ix < EI_CLASSIFIER_LABEL_COUNT; ix++) {
114             ei_printf(" %s: %.5f\n", result.classification[ix].label,
115                 result.classification[ix].value);
116         }
117         if (EI_CLASSIFIER_HAS_ANOMALY == 1)
118             ei_printf(" anomaly score: %.5f\n", result.anomaly);
119         print_results = 0;
120     }
121 }
122 }

```

FIGURE 5: Code first half

```

1
2
3 /**      Printf function uses vsprintf and output using Arduino Serial
4  *param[in] format      Variable argument list
5  */
6 void ei_printf(const char *format, ...) {
7     static char print_buf[1024] = { 0 };
8
9     va_list args;
10    va_start(args, format);
11    int r = vsprintf(print_buf, sizeof(print_buf), format, args);
12    va_end(args);
13
14    if (r > 0) {
15        Serial.write(print_buf);
16    }
17 }
18
19 /**
20  *brief      PDM buffer full callback
21  *          Get data and call audio thread callback
22  */
23 static void pdm_data_ready_inference_callback(void)
24 {
25     int bytesAvailable = PDM.available();
26
27     // read into the sample buffer
28     int bytesRead = PDM.read((char *)&sampleBuffer[0], bytesAvailable);
29
30     if (record_ready == true) {
31         for (int i = 0; i < bytesRead; i++) {
32             inference_buffers[inference_buf_select][inference_buf_count++] = sampleBuffer[i];
33         }
34         if (inference_buf_count == inference_n_samples) {
35             inference_buf_select++;
36             inference_buf_count = 0;
37             inference_buf_ready = 1;
38         }
39     }
40 }
41
42 /**
43  *brief      Init inferencing struct and setup start PDM
44  *param[in] n_samples The n samples
45  *return      description of the return value
46  */
47 static bool microphone_inference_start(uint32_t n_samples)
48 {
49     inference_buffers[0] = (signed short *)malloc(n_samples * sizeof(signed short));
50
51     if (inference_buffers[0] == NULL) {
52         return false;
53     }
54
55     inference_buffers[1] = (signed short *)malloc(n_samples * sizeof(signed short));
56
57     if (inference_buffers[0] == NULL) {
58         free(inference_buffers[0]);
59         return false;
60     }
61
62     sampleBuffer = (signed short *)malloc((n_samples > 1) * sizeof(signed short));
63
64     if (sampleBuffer == NULL) {
65         free(inference_buffers[0]);
66         free(inference_buffers[1]);
67         return false;
68     }
69
70     inference_buf_select = 0;
71     inference_buf_count = 0;
72     inference_n_samples = n_samples;
73     inference_buf_ready = 0;
74
75     // configure the data receive callback
76     PDM.onReceive(&pdm_data_ready_inference_callback);
77
78     PDM.setBufferSize((n_samples > 1) * sizeof(int16_t));
79
80     // Initialize PDM with:
81     // - one channel (mono mode)
82     // - a 16 kHz sample rate
83     if (!PDM.begin(1, EI_CLASSIFIER_FREQUENCY)) {
84         ei_printf("failed to start PDM");
85     }
86
87     // set the gain, defaults to 20
88     PDM.setGain(20);
89
90     record_ready = true;
91     return true;
92 }
93
94 /**
95  *brief      Wait on new data
96  *return      True when finished
97  */
98 static bool microphone_inference_record(void)
99 {
100     bool ret = true;
101
102     if (inference_buf_ready == 1) {
103         ei_printf(
104             "Error sample buffer overrun. Decrease the number of slices per model window "
105             "(EI_CLASSIFIER_SLICES_PER_MODEL_WINDOW)\n");
106         ret = false;
107     }
108     while (inference_buf_ready == 0) {
109         delay(1);
110     }
111     inference_buf_ready = 0;
112     return ret;
113 }
114
115 /** Get raw audio signal data
116  */
117 static int microphone_audio_signal_get_data(size_t offset, size_t length, float *out_ptr)
118 {
119     numpy::int16_to_float(inference_buffers[inference_buf_select + 1][offset], out_ptr, length);
120     return 0;
121 }
122
123 /** Stop PDM and release buffers
124  */
125 static void microphone_inference_end(void)
126 {
127     PDM.end();
128     free(inference_buffers[0]);
129     free(inference_buffers[1]);
130     free(sampleBuffer);
131 }
132
133 #if !defined(EI_CLASSIFIER_SENSOR) || EI_CLASSIFIER_SENSOR != EI_CLASSIFIER_SENSOR_MICROPHONE
134 #error "Invalid model for current sensor."
135 #endif

```

FIGURE 6: Code second half